

part 2: scenarios:

SCENARIO 1: Two Prophets, One Game

Average test error: 0.3067
 Approximation error: 0.2000
 Estimation error: 0.0940
 Best prophet chosen: 53/100 times

SCENARIO 2: Two Prophets, Ten Games

Average test error: 0.2213
 Approximation error: 0.2000
 Estimation error: 0.0240
 Best prophet chosen: 88/100 times

SCENARIO 3: Many Prophets, Ten Games

Average test error: 0.0930
 Approximation error: 0.0064
 Estimation error: 0.0852
 Best prophet chosen: 3/100 times
 Within 1% of best: 8/100 times

analysis: if the error rates were uniformly distributed between [0, 0.5] instead of [0, 1], the **approximation error would stay largely the same**, since we have haven't changes the lower bound on True risk. meanwhile, the **estimation error would be substantially smaller**, since even when we select a model whose true risk is greater than that of the best available model, the true risk of the chosen model will be closer to 0.

SCENARIO 4: Many Prophets, Many Games

Average test error: 0.0064
 Approximation error: 0.0064
 Estimation error: 0.0008
 Best prophet chosen: 50/100 times
 Within 1% of best: 98/100 times

analysis: if we evaluate the generalisation gap of a model based on the train set, we expect it to be **greater** than if we had measured generalisation gap based on the model's performance on the test. That is because the train set is a relatively small subset compared to the population and the model selection was biased in favor of a model performing well on the train set. No such bias affects the performance on the test set, and hence the test set provides a better approximation of the generalisation gap.

SCENARIO 5: School of Prophets

Grid search completed

k values (prophets): [2, 5, 10, 50]
m values (train games): [1, 10, 50, 1000]
Number of trials: 100

Scenario 5: School of Prophets Results														
Average Test Error				Approximation Error				Estimation Error						
	m=1	m=10	m=50	m=1000		m=1	m=10	m=50	m=1000		m=1	m=10	m=50	m=1000
k=2	0.0987	0.0861	0.0738	0.0703	k=2	0.0643	0.0730	0.0683	0.0707	k=2	0.0341	0.0126	0.0048	0.0003
k=5	0.1006	0.0658	0.0409	0.0294	k=5	0.0346	0.0340	0.0351	0.0298	k=5	0.0667	0.0332	0.0053	0.0002
k=10	0.1019	0.0665	0.0239	0.0187	k=10	0.0184	0.0176	0.0155	0.0182	k=10	0.0836	0.0486	0.0083	0.0003
k=50	0.0960	0.0653	0.0178	0.0034	k=50	0.0036	0.0036	0.0034	0.0030	k=50	0.0925	0.0623	0.0143	0.0004

The grid search shows clear patterns: as we increase **k** (class size) we increase two factors:

1. The chance of the available prophets including a prophet with a lower true risk, thus **decreasing approximation error**. Indeed **k** is almost entirely responsible for the approximation error.
2. Increasing the chance of including a prophet that happened to perform well on the train set, while having a greater true risk, thus **increasing estimation error and test error**. Meanwhile, increasing **m** (the training set size) makes it harder for a model to perform well on the train set despite having a higher true risk. Thus, it **decreases the estimation error and the test error**.

Indeed, we see that when we increased the number of games for ERM from scenario 1 to 2, we lowered the estimation and test error as expected.

Furthermore, scenarios 3 and 4 exactly show how the approximation error is affected primarily by **k**, while increasing **m** led to lower test and estimation errors.

SCENARIO 6: Bias-Complexity Tradeoff

Hypothesis 1 (5 prophets, high bias):

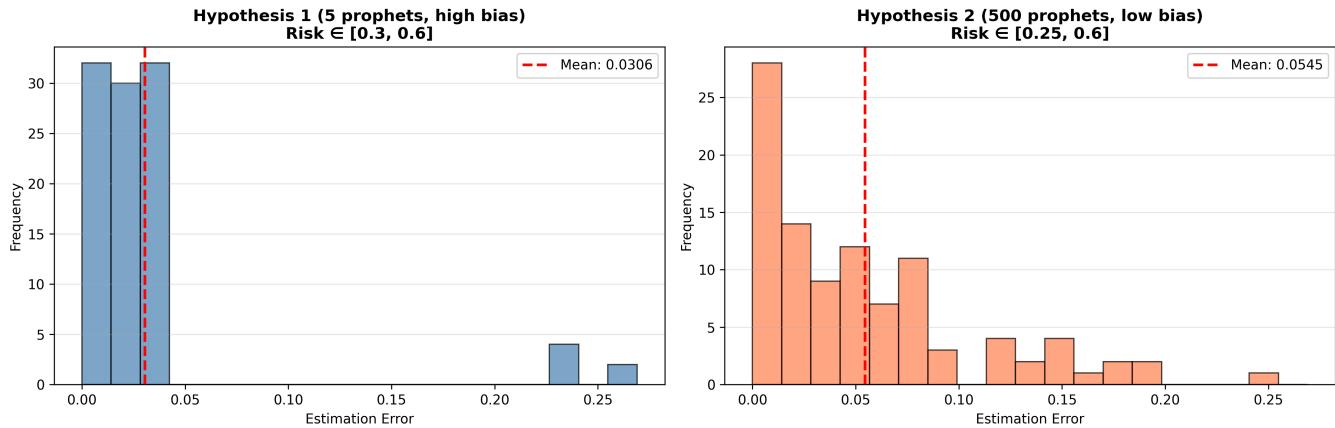
Average test error: 0.3570

Approximation error: 0.3230
 Estimation error: 0.0306

Hypothesis 2 (500 prophets, low bias):

Average test error: 0.3044
 Approximation error: 0.2500
 Estimation error: 0.0545

Scenario 6: Bias-Complexity Tradeoff - Estimation Error Distribution



we see that while class 2 has models sampled from a 'better' range of true risk (lower bound 0.25 as opposed to 0.3), the estimation error is higher. this is as expected, since given a larger class size there is a greater likelihood of choosing a sample with low error on the train set despite higher true risk. meanwhile, the fact that the models in class 2 on average have lower true risk, leads to a lower test error.

part 3: pac learning analysis:

in this part we use the formula relating the number of samples to accuracy and confidence: $m = 2 * np.log(2 * h / \delta) / \epsilon$, which can also be expressed as: $m = (2 / \epsilon) * (np.log(2) + np.log(h) - np.log(\delta))$ we note that m is inversely proportional to epsilon and to the log of delta. it is proportional to the log of h.

QUESTION 1: Compute minimal number of samples

Input Parameters:

$|H|$ (hypothesis class size) = 100
 ϵ (desired accuracy) = 0.05
 δ (confidence level) = 0.01

Output: m (rounded up) = 397

QUESTION 2: Analyze change when $|H|$ is doubled

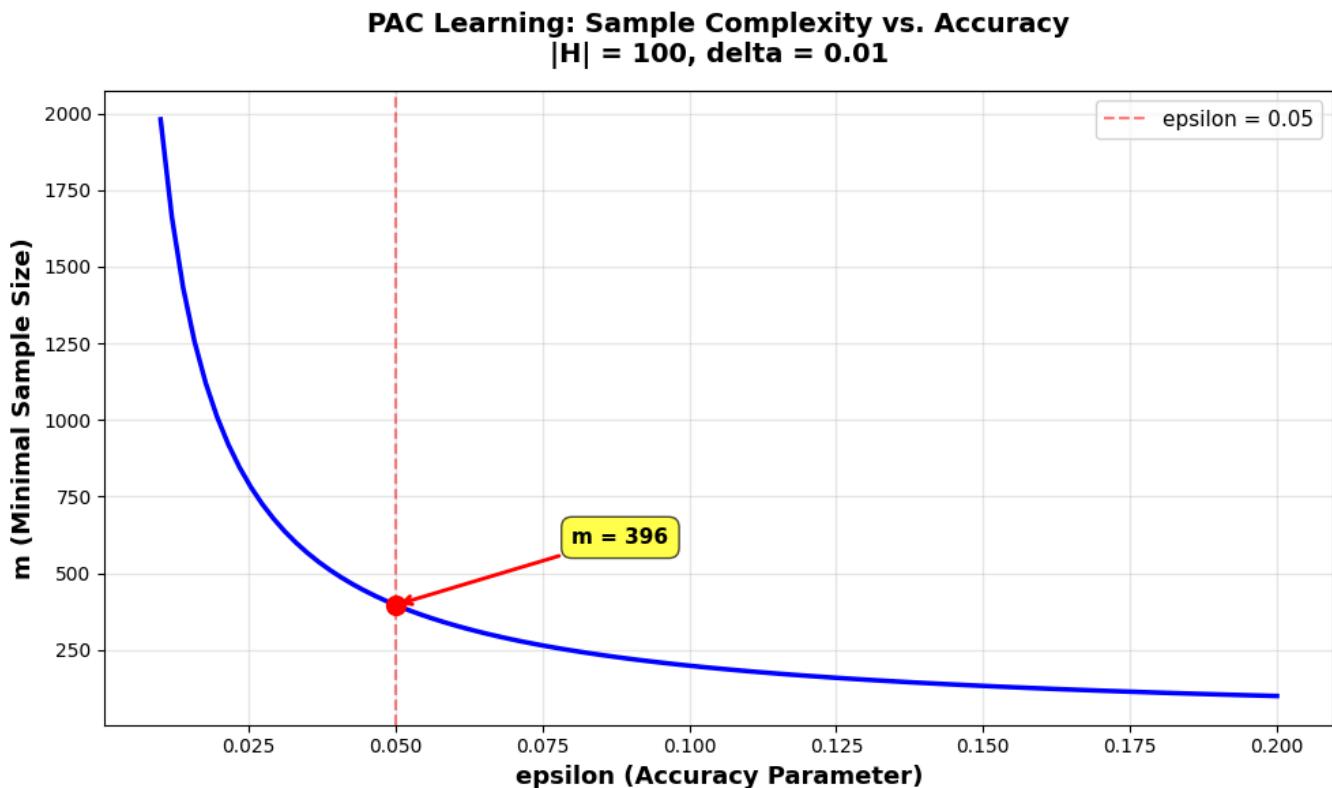
Input Parameters:

$|H|$ (hypothesis class size) = 200 (doubled)

Output: m (rounded up) = 424

we see that if we double $|H|$ we only need an additional 27 samples. if we refactor the equation we get ' $m = \text{original_val} + (2 / \text{epsilon}) * \text{np.log}(2)$ ' which comes to 27

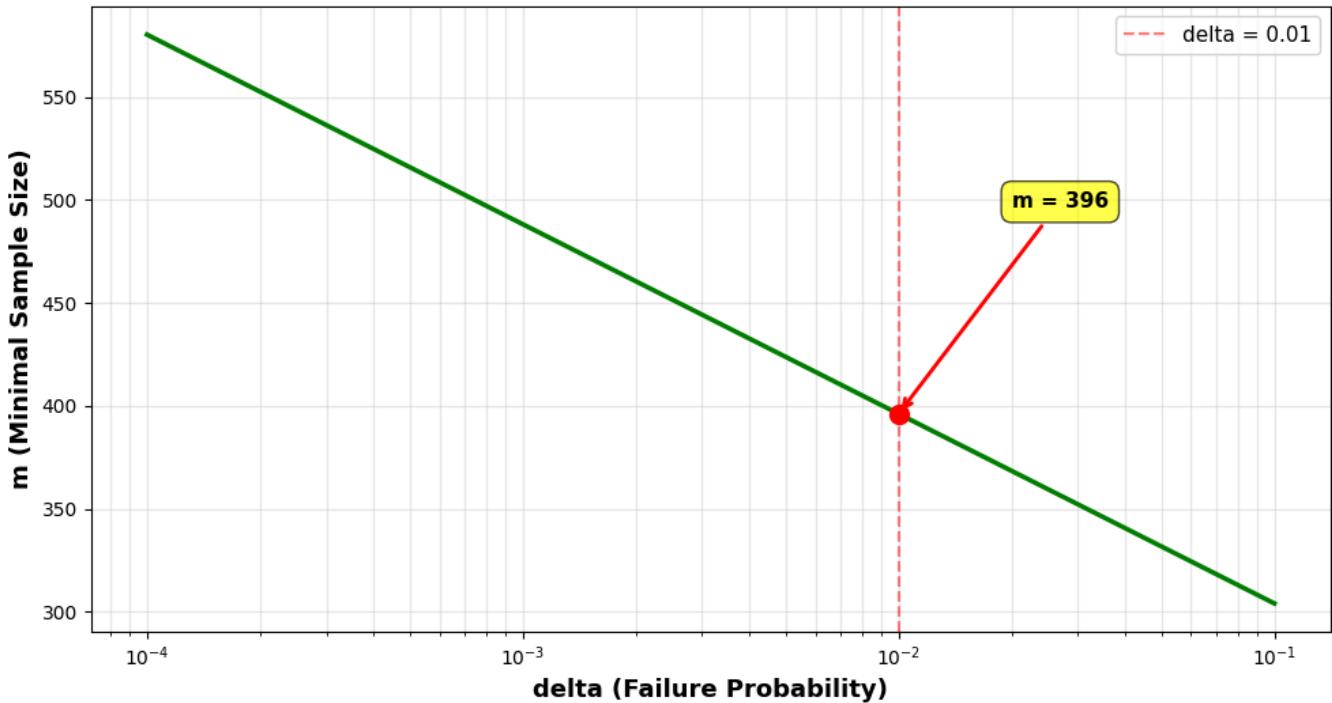
QUESTION 3: Plot m as a function of epsilon in $[0.01, 0.2]$



we clearly see m increase in inverse proportion to epsilon. hence, to halve the error, we need to double the number of samples.

QUESTION 4: Plot m as a function of delta in $[10^{-4}, 0.1]$ (log scale)

PAC Learning: Sample Complexity vs. Confidence
 $|H| = 100$, $\epsilon = 0.05$



we clearly see that curve follows an exact logarithmic scale. hence, to increase the confidence 10 fold, we need only increase the numerator by a factor of $\log(10)$, making this scale well.