

## 1 Part A: Calculus

### 1.1 Calculus-Q1

CQ1.1 Use the chain rule to calculate the gradient of

$$h(\sigma) = \frac{1}{2} \|f(\sigma) - y\|^3$$

where  $\sigma \in \mathbb{R}^m$  and  $f$  is some arbitrary function from  $\mathbb{R}^m$  to  $\mathbb{R}^n$ .

CQ1.2 Compute the expression in the case where:

$$f(\sigma) = \begin{bmatrix} \sigma_1 \cdot \sigma_2 \\ \sigma_1^2 + \sigma_2^2 \\ \sigma_1 \end{bmatrix}$$

$$\sigma = \begin{bmatrix} 1 \\ 2 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

$$h(\sigma) = \frac{1}{2} \|r(\sigma)\|^3$$

let:

$$r(\sigma) = f(\sigma) - y, \quad g(u) = \frac{1}{2} \|u\|^3$$

hence:

$$h(\sigma) = g(r(\sigma))$$

Let us find gradients:

$$\|u\| = (u \cdot u^\top)^{\frac{1}{2}}$$

$$g(u) = \frac{1}{2} \|u\|^3 = \frac{1}{2} (u \cdot u^\top)^{\frac{3}{2}}$$

$$g'(u) = \frac{1}{2} \cdot \frac{3}{2} \cdot (u \cdot u^\top)^{\frac{1}{2}} \cdot (u \cdot u^\top)' = \frac{1}{2} \cdot \frac{3}{2} \cdot (u \cdot u^\top)^{\frac{1}{2}} \cdot (u^2)' = \frac{1}{2} \cdot \frac{3}{2} \cdot \|u\| \cdot 2u = \frac{3}{2} \|u\| \cdot u =$$

let  $u = r(\sigma)$ :

$$g'(u) = \frac{3}{2} \|r(\sigma)\| r(\sigma)$$

we evaluate  $r(\sigma)$   
as the jacobian:

$$J_f(\sigma) = \begin{bmatrix} \frac{\partial f_1}{\partial \sigma_1} & \dots & \frac{\partial f_1}{\partial \sigma_m} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_n}{\partial \sigma_1} & \dots & \frac{\partial f_n}{\partial \sigma_m} \end{bmatrix} \in \mathbb{R}^{n \times m}$$

$$r(\sigma) = \begin{pmatrix} 2 & 1 \\ 2 & 4 \\ 1 & 0 \end{pmatrix}$$

we want to find gradient  $h(\sigma)$

evaluated at:

$$f(\sigma) = \begin{bmatrix} \sigma_1 \cdot \sigma_2 \\ \sigma_1^2 + \sigma_2^2 \\ \sigma_1 \end{bmatrix}$$

$$\sigma = \begin{bmatrix} 1 \\ 2 \end{bmatrix}, \quad y = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

by the chain rule:

$$h'(\sigma) = g'(r(\sigma)) \cdot r'(\sigma)$$

$$= \frac{3}{2} \|r(\sigma)\| \cdot r(\sigma) \cdot f'(\sigma)$$

$$= \frac{3}{2} \|r(\sigma)\| \cdot J_f(\sigma)^T \cdot r(\sigma)$$

let us evaluate:

$$\sigma = \begin{bmatrix} 1 \\ 2 \end{bmatrix} \quad y = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$$

$$f(\sigma) = \begin{pmatrix} 2 \\ 5 \\ 1 \end{pmatrix} \Rightarrow f'(\sigma) = \begin{pmatrix} 2 & 1 \\ 2 & 4 \\ 1 & 0 \end{pmatrix}$$

$$r(\sigma) = \begin{pmatrix} 1 \\ 4 \\ 0 \end{pmatrix}$$

$$h'(\sigma) = \frac{3}{2} \cdot \| \begin{pmatrix} 1 \\ 4 \\ 0 \end{pmatrix} \| \cdot \begin{pmatrix} 2 & 1 \\ 2 & 4 \\ 1 & 0 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 4 \\ 0 \end{pmatrix}$$

$$\boxed{\frac{3}{2} \cdot \sqrt{17} \cdot \begin{pmatrix} 10 \\ 17 \end{pmatrix}}$$

## 1.2 Calculus-Q2

The softmax function  $S : \mathbb{R}^k \rightarrow [0, 1]^k$ , is defined as follows:

$$S(x)_j = \frac{e^{x_j}}{\sum_{l=1}^k e^{x_l}}$$

This function takes an input vector  $x \in \mathbb{R}^d$  and outputs a probability vector (non-negative entries that sum up to 1), corresponding to the weight of original entries of  $x$ .

CQ2: Calculate the Jacobian of the softmax function  $S$ .

$$S'(x) = \begin{bmatrix} \frac{\partial S_1}{\partial x_1} & \dots & \frac{\partial S_1}{\partial x_K} \\ \vdots & & \vdots \\ \frac{\partial S_K}{\partial x_1} & \dots & \frac{\partial S_K}{\partial x_K} \end{bmatrix}$$

$$\text{let } Z = \sum_{i=1}^K e^{x_i}.$$

$$\begin{aligned} \frac{\partial S_i}{\partial x_j} &:= \frac{1}{Z} \cdot \frac{\cancel{e^{x_i}}}{\cancel{e^{x_j}}} * \frac{\cancel{\frac{d(e^{x_i})}{dx_j}} \cdot Z - e^{x_i} \cdot \cancel{\frac{d(Z)}{dx_j}}}{Z^2} \\ &\stackrel{**}{=} \frac{S_{ij} e^{x_i} Z - e^{x_i} e^{x_j}}{Z^2} \\ &= \frac{e^{x_i}}{Z} \left( S_{ij} - \frac{e^{x_j}}{Z} \right) \end{aligned}$$

therefore

$$J_{ij} = \begin{cases} S(x)_i \cdot (1 - S(x)) & i=j \\ -\frac{e^{x_i} e^{x_j}}{Z^2} = -S(x)_i S(x)_j & i \neq j \end{cases}$$

$$\begin{aligned} * \frac{f(x)}{g(x)} &= \frac{f'(x)g(x) - f(x)g'(x)}{g(x)^2} \\ (\sum e^{x_i})' &= \sum (e^{x_i})' = \sum e^{x_i} \\ ** \frac{\frac{d}{dx_j} e^{x_i}}{Z} &= \begin{cases} e^{x_i} & i=j \\ 0 & i \neq j \end{cases} \\ S_{ij} &= \begin{cases} 1 & i=j \\ 0 & i \neq j \end{cases} \\ \frac{d}{dx_j}(Z) &= \frac{1}{Z} \sum e^{x_i} = \begin{cases} 1 & i=j \\ 0 & i \neq j \end{cases} \end{aligned}$$

$$f = \begin{bmatrix} S(x)_1 \cdot (1 - S(x)_1) & -S(x)_1 \cdot S(x)_2 & \dots & -S(x)_1 \cdot S(x_K) \\ \vdots & \ddots & S(x_2) \cdot (1 - S(x_2)) & \\ -S(x_K) \cdot S(x_1) & & \ddots & S(x_K) \cdot (1 - S(x_K)) \end{bmatrix}$$

## part 2: scenarios:

---

### SCENARIO 1: Two Prophets, One Game

Average test error:	0.3067
Approximation error:	0.2000
Estimation error:	0.0940
Best prophet chosen:	53/100 times

### SCENARIO 2: Two Prophets, Ten Games

Average test error:	0.2213
Approximation error:	0.2000
Estimation error:	0.0240
Best prophet chosen:	88/100 times

### SCENARIO 3: Many Prophets, Ten Games

Average test error:	0.0930
Approximation error:	0.0064
Estimation error:	0.0852
Best prophet chosen:	3/100 times
Within 1% of best:	8/100 times

analysis: if the error rates were uniformly distributed between [0, 0.5] instead of [0, 1], the **approximation error would stay largely the same**, since we have haven't changes the lower bound on True risk. meanwhile, the **estimation error would be substantially smaller**, since even when we select a model whose true risk is greater than that of the best available model, the true risk of the chosen model will be closer to 0.

### SCENARIO 4: Many Prophets, Many Games

Average test error:	0.0064
Approximation error:	0.0064
Estimation error:	0.0008
Best prophet chosen:	50/100 times
Within 1% of best:	98/100 times

analysis: if we evaluate the generalisation gap of a model based on the train set, we expect it to be **greater** than if we had measured generalisation gap based on the model's performance on the test. That is because the train set is a relatively small subset compared to the population and the model selection was biased in favor of a model performing well on the train set. No such bias affects the performance on the test set, and hence the test set provides a better approximation of the generalisation gap.

## SCENARIO 5: School of Prophets

Grid search completed

k values (prophets): [2, 5, 10, 50]  
m values (train games): [1, 10, 50, 1000]  
Number of trials: 100

Scenario 5: School of Prophets Results														
Average Test Error				Approximation Error				Estimation Error						
	m=1	m=10	m=50	m=1000		m=1	m=10	m=50	m=1000		m=1	m=10	m=50	m=1000
k=2	0.0987	0.0861	0.0738	0.0703	k=2	0.0643	0.0730	0.0683	0.0707	k=2	0.0341	0.0126	0.0048	0.0003
k=5	0.1006	0.0658	0.0409	0.0294	k=5	0.0346	0.0340	0.0351	0.0298	k=5	0.0667	0.0332	0.0053	0.0002
k=10	0.1019	0.0665	0.0239	0.0187	k=10	0.0184	0.0176	0.0155	0.0182	k=10	0.0836	0.0486	0.0083	0.0003
k=50	0.0960	0.0653	0.0178	0.0034	k=50	0.0036	0.0036	0.0034	0.0030	k=50	0.0925	0.0623	0.0143	0.0004

The grid search shows clear patterns: as we increase **k** (class size) we increase two factors:

1. The chance of the available prophets including a prophet with a lower true risk, thus **decreasing approximation error**. Indeed **k** is almost entirely responsible for the approximation error.
2. Increasing the chance of including a prophet that happened to perform well on the train set, while having a greater true risk, thus **increasing estimation error and test error**. Meanwhile, increasing **m** (the training set size) makes it harder for a model to perform well on the train set despite having a higher true risk, thus **decreases the estimation error and the test error**.

Indeed, we see that when we increased the number of games for ERM from scenario 1 to 2, we lowered the estimation and test error as expected.

Furthermore, scenarios 3 and 4 exactly show how the approximation error is affected primarily by **k**, while increasing **m** led to lower test and estimation errors.

## SCENARIO 6: Bias-Complexity Tradeoff

Hypothesis 1 (5 prophets, high bias):

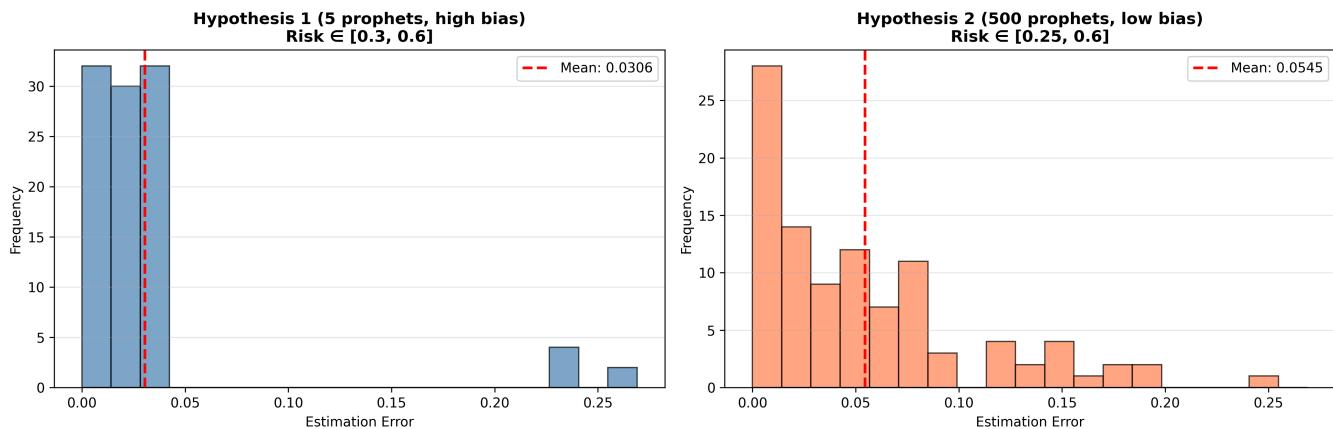
Average test error: 0.3570

Approximation error: 0.3230  
 Estimation error: 0.0306

Hypothesis 2 (500 prophets, low bias):

Average test error: 0.3044  
 Approximation error: 0.2500  
 Estimation error: 0.0545

**Scenario 6: Bias-Complexity Tradeoff - Estimation Error Distribution**



we see that while class 2 has models sampled from a 'better' range of true risk (lower bound 0.25 as opposed to 0.3), the estimation error is higher. this is as expected, since given a larger class size there is a greater likelihood of choosing a sample with low error on the train set despite higher true risk. meanwhile, the fact that the models in class 2 on average have lower true risk, leads to a lower test error.

## part 3: pac learning analysis:

in this part we use the formula relating the number of samples to accuracy and confidence:  $m = 2 * np.log(2 * h / \delta) / \epsilon$ , which can also be expressed as:  $m = (2 / \epsilon) * (np.log(2) + np.log(h) - np.log(\delta))$  we note that m is inversely proportional to epsilon and to the log of delta. it is proportional to the log of h.

### QUESTION 1: Compute minimal number of samples

Input Parameters:

$|H|$  (hypothesis class size) = 100  
 $\epsilon$  (desired accuracy) = 0.05  
 $\delta$  (confidence level) = 0.01

Output: m (rounded up) = 397

### QUESTION 2: Analyze change when $|H|$ is doubled

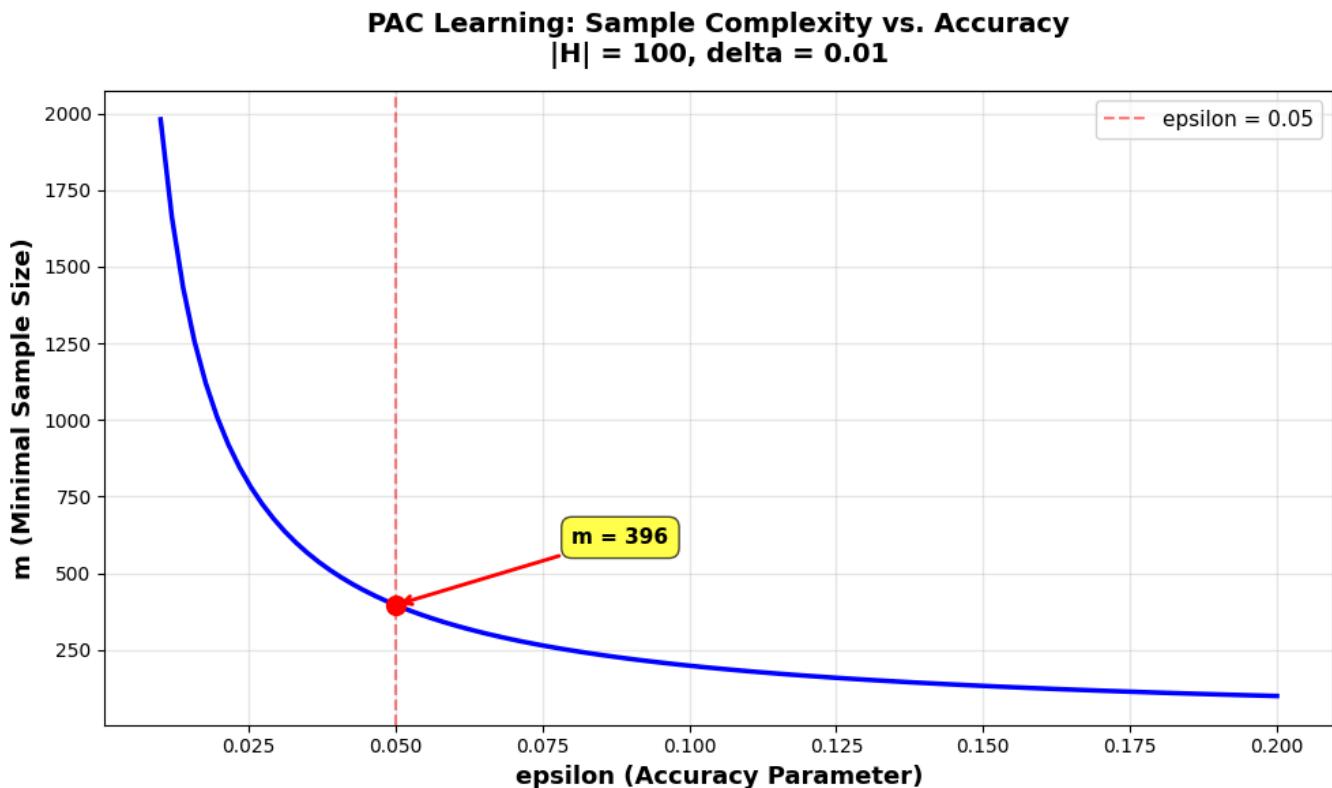
**Input Parameters:**

$|H|$  (hypothesis class size) = 200 (doubled)

Output:  $m$  (rounded up) = 424

we see that if we double  $|H|$  we only need an additional 27 samples. if we refactor the equation we get ' $m = \text{original\_val} + (2 / \text{epsilon}) * \text{np.log}(2)$ ' which comes to 27

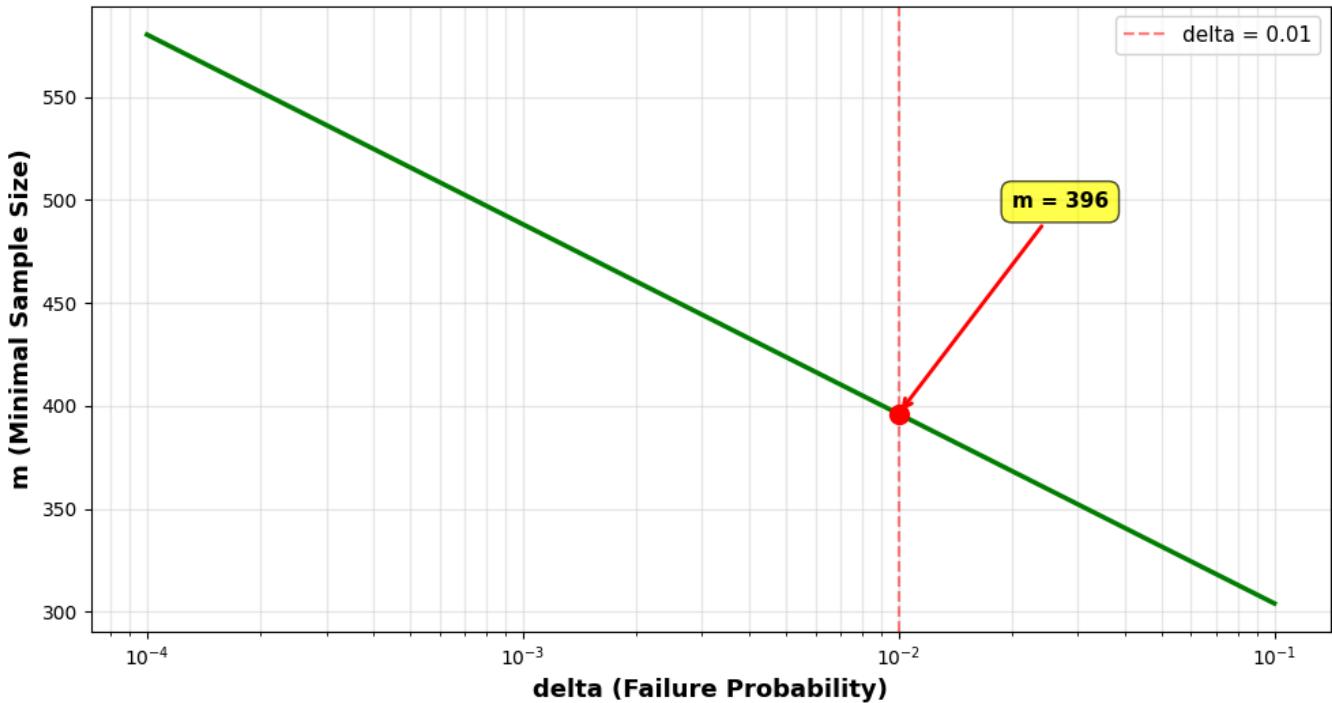
**QUESTION 3:** Plot  $m$  as a function of epsilon in  $[0.01, 0.2]$



we clearly see  $m$  increase in inverse proportion to epsilon. hence, to halve the error, we need to double the number of samples.

**QUESTION 4:** Plot  $m$  as a function of delta in  $[10^{-4}, 0.1]$  (log scale)

**PAC Learning: Sample Complexity vs. Confidence**  
 $|H| = 100$ ,  $\epsilon = 0.05$



we clearly see that curve follows an exact logarithmic scale. hence, to increase the confidence 10 fold, we need only increase the numerator by a factor of  $\log(10)$ , making this scale well.